

Exam Revision

Edward Zhang

SOFTENG 370 T8

Exam Info

Your exam will be short answer, not MCQ. That means the exam from 2012 - 2017 aren't very useful. 2018 had a different lecturer for the first $\frac{1}{4}$ so it's not super helpful either.

Which of the following is not a necessary component of a monitor?

- ▶ Publicly accessible entry points
- ▶ A readers/writers lock
- ▶ A scheduler
- ▶ A shared resource which is protected by the monitor

Which of the following is not a necessary component of a monitor?

- ▶ Publicly accessible entry points
- ▶ A readers/writers lock
- ▶ A scheduler
- ▶ A shared resource which is protected by the monitor

Explanation

Reader/Writers lock can enhance performance, but is not required.

Which of the following best explains what happens when a damaged C program comes to an end but doesn't call the exit routine?

- ▶ The damaged program can corrupt memory used by other processes and cause them to crash or perform illegal instructions.
- ▶ The operating system takes control when the program tries to execute an illegal instruction or attempts to access unallocated memory.
- ▶ The C standard library takes control when the program fails to return to the code which called the main function.
- ▶ The operating system creates a new process and restarts the damaged program in that process so that it gets another chance to complete.

Which of the following best explains what happens when a damaged C program comes to an end but doesn't call the exit routine?

- ▶ The damaged program can corrupt memory used by other processes and cause them to crash or perform illegal instructions.
- ▶ The operating system takes control when the program tries to execute an illegal instruction or attempts to access unallocated memory.
- ▶ The C standard library takes control when the program fails to return to the code which called the main function.
- ▶ The operating system creates a new process and restarts the damaged program in that process so that it gets another chance to complete.

The code below uses a compare and swap function “cas”. What is the code doing?

```
add_to_balance(increase):  
    previous_amount = balance  
    while (!cas(&balance,  
               previous_amount,  
               previous_amount + increase)):  
        previous_amount = balance
```

- ▶ It repeatedly increments balance by increase until balance overflows.
- ▶ It increments balance by increase using a condition variable.
- ▶ It safely swaps the values of balance with previous_amount + balance using a wait-free algorithm.
- ▶ It safely increments balance by increase using a lock-free algorithm.

The code below uses a compare and swap function “cas”. What is the code doing?

```
add_to_balance(increase):  
    previous_amount = balance  
    while (!cas(&balance,  
               previous_amount,  
               previous_amount + increase)):  
        previous_amount = balance
```

- ▶ It repeatedly increments balance by increase until balance overflows.
- ▶ It increments balance by increase using a condition variable.
- ▶ It safely swaps the values of balance with previous_amount + balance using a wait-free algorithm.
- ▶ It safely increments balance by increase using a lock-free algorithm.

Which of the following does NOT happen in a context switch between threads in the same process?

- ▶ The processor registers for the currently running thread are saved.
- ▶ The processor registers are loaded with the saved values for the new thread.
- ▶ The page table is switched from the old thread to the new thread.
- ▶ The thread states for the two threads may be changed.
- ▶ The stack is changed from the old thread to the new thread.

Which of the following does NOT happen in a context switch between threads in the same process?

- ▶ The processor registers for the currently running thread are saved.
- ▶ The processor registers are loaded with the saved values for the new thread.
- ▶ The page table is switched from the old thread to the new thread.
- ▶ The thread states for the two threads may be changed.
- ▶ The stack is changed from the old thread to the new thread.

Explanation

Memory is shared between threads, so same page table.

Which of the following is False?

- ▶ FUSE works by redirecting file operations through the FUSE module to a process running in user mode.
- ▶ To use a FUSE file system we mount the file system over an existing directory.
- ▶ To use FUSE requires root privileges.
- ▶ If the FUSE process is killed the files and directories contained within it will not be accessible.
- ▶ There has to be a FUSE kernel module in order for FUSE to work on Linux.

Which of the following is False?

- ▶ FUSE works by redirecting file operations through the FUSE module to a process running in user mode.
- ▶ To use a FUSE file system we mount the file system over an existing directory.
- ▶ To use FUSE requires root privileges.
- ▶ If the FUSE process is killed the files and directories contained within it will not be accessible.
- ▶ There has to be a FUSE kernel module in order for FUSE to work on Linux.

Explanation

You probably used FUSE w/o root in your assignment.

Which of the following disk scheduling algorithms are commonly used for scheduling SSDs?

- ▶ Shortest Seek time First
- ▶ SCAN
- ▶ First come, first served
- ▶ Circular SCAN
- ▶ None of the above

Which of the following disk scheduling algorithms are commonly used for scheduling SSDs?

- ▶ Shortest Seek time First
- ▶ SCAN
- ▶ **First come, first served**
- ▶ Circular SCAN
- ▶ None of the above

Explanation

SSDs have no Seek time, and no head/platter so SCAN is irrelevant. FCFS makes sense since no special handling is required.

What causes thrashing?

- ▶ When the foreground process has completely used up the number of frames it has been allocated.
- ▶ When the sum of the pages of the working-sets exceeds the number of frames.
- ▶ When there is not enough contiguous memory to be allocated for all current working sets.
- ▶ When all frames are currently being used.
- ▶ When all processes have filled up their page tables.

What causes thrashing?

- ▶ When the foreground process has completely used up the number of frames it has been allocated.
- ▶ When the sum of the pages of the working-sets exceeds the number of frames.
- ▶ When there is not enough contiguous memory to be allocated for all current working sets.
- ▶ When all frames are currently being used.
- ▶ When all processes have filled up their page tables.

Explanation

Recall that thrashing is when the virtual memory system is overused, and is thus stuck in a constant state of paging / pagefaults.

Which of the following statements about user level device drivers is FALSE?

- ▶ User level drivers cannot deal with device interrupts.
- ▶ User level drivers can communicate with memory mapped devices.
- ▶ Most problems with user level drivers do not affect the kernel.
- ▶ Because of mode transitions user level drivers are sometimes not used for fast devices.
- ▶ User level drivers can communicate with IO ports

Which of the following statements about user level device drivers is FALSE?

- ▶ User level drivers cannot deal with device interrupts.
- ▶ User level drivers can communicate with memory mapped devices.
- ▶ Most problems with user level drivers do not affect the kernel.
- ▶ Because of mode transitions user level drivers are sometimes not used for fast devices.
- ▶ User level drivers can communicate with IO ports

Question 9b

One protection against the Meltdown exploit that has been implemented in operating systems is kernel page-table isolation (KPTI). Explain what kernel page-table isolation is.

Question 9b

One protection against the Meltdown exploit that has been implemented in operating systems is kernel page-table isolation (KPTI). Explain what kernel page-table isolation is.

Answer

Separate page tables are kept for a process when it is running in kernel and user mode. The user mode page tables do not have most of the kernel pages mapped into them.

Question 9c/d

What effect could KPTI have on efficiency, and how does it prevent Meltdown?

Question 9c/d

What effect could KPTI have on efficiency, and how does it prevent Meltdown?

Efficiency

Anything that switches to kernel mode, like syscalls, requires a new page table to be loaded. This will also likely flush the TLB.

Question 9c/d

What effect could KPTI have on efficiency, and how does it prevent Meltdown?

Efficiency

Anything that switches to kernel mode, like syscalls, requires a new page table to be loaded. This will also likely flush the TLB.

Preventing Meltdown

Meltdown's timing attack relies on attempts to access values in kernel address space (even if these are eventually rejected by a privilege check). However, with KPTI, those addresses are not accessible in the current address space, so the attempt cannot be made.